

SOURCE OF TRUTH

Centro Architecture & Philosophy

Pilot Edition — Revision 21 (Landing Page — Contact Form Email & WhatsApp Button)

This document defines the architectural principles and product philosophy that guide every design and development decision in Centro. It is version-controlled in the repository as `ARCHITECTURE.md` and rendered to PDF after every roadmap milestone.

Status	Living document — source of truth, updated at the end of every milestone
Governs	Every design and implementation decision in this codebase
Revision history	See "Document History" at the end

Chapter 1 — Learning First

Every client begins in Learning Mode.

Centro does not require a perfect setup before work begins. Onboarding exists to provide enough information to start working safely, not to perfectly configure every client.

Excel imports, onboarding choices, and existing records create only an Initial Client Profile. From the very first collection cycle, Centro observes reality, learns from recurring behaviour, and continuously improves the profile.

Centro does not guess. It observes, suggests, confirms, and only then learns.

Core principles:

- Start working immediately.
- Learn continuously.
- Never guess when confidence is low.
- Improve from real behaviour, not assumptions.

Chapter 2 — The First Collection Cycle

Every client enters the first month in Learning Mode.

Centro requests the client's documents, waits until the client confirms that the submission is complete, and then treats the collection as one business event to learn from.

The first cycle is not a validation phase. It is a learning phase designed to understand how the client actually works — specifically, which documents that client actually sends, nothing broader.

Chapter 3 — Observe → Suggest → Confirm → Learn

Every meaningful change to a client's document collection profile follows the same lifecycle:

1. **Observe.**
2. **Suggest.**
3. **Confirm** with the client whenever appropriate, through WhatsApp.
4. **Learn.**

Centro never changes a client's profile because of a single event or assumption. A pattern must repeat before Centro treats it as worth asking about, and even then, only the client's explicit confirmation — never the pattern alone — changes the profile.

Chapter 4 — Learning Without Becoming a Burden

Learning must never create unnecessary work.

Centro should quietly observe whenever possible. It asks questions only if the answer has long-term value and the benefit outweighs the interruption.

Questions must be short, natural, and asked only once — if a client has already declined a suggestion, Centro does not ask about the same thing again.

The accountant reviews exceptions — not routine learning.

Core principle: Centro becomes smarter while becoming quieter.

Chapter 5 — The Living Client Profile

Client profiles are living profiles.

Knowledge comes from Excel, onboarding, collected documents, recurring patterns, and confirmed client responses.

No single source is absolute truth.

Businesses evolve, so the client's document collection profile evolves with them — see Chapter 8 for the exact, permanent boundary of what "evolves" is allowed to mean.

Chapter 6 — Decision Hierarchy

Every decision follows the same hierarchy, most specific and most certain signal first:

1. **Learned Knowledge** — a specific fact this organization has already confirmed about its own clients or data. Checked first because it is the most reliable signal available: it was explicitly confirmed, not inferred, and it is scoped to exactly this organization.
2. **Business Rules** — deterministic, universal rules that apply the same way across every organization (e.g. recognizing standard terminology, hard confidence floors, required-confirmation gates). Used whenever no more specific learned fact exists for this organization yet.
3. **Artificial Intelligence** — a supporting tool, reached only once the above two layers find nothing with sufficient confidence. AI is never the foundation of a decision in Centro.
4. **Human Review** — whenever confidence remains insufficient after all three layers above, a person decides. Centro never guesses.

A previously confirmed, organization-specific fact always outranks a generic rule. This is a deliberate refinement of "business rules first": a rule that is merely generic is less trustworthy than a fact Centro was explicitly told and is now honoring. Business Rules remain the universal floor beneath Learned Knowledge — every organization gets at least this level of correctness even before it has taught Centro anything of its own — but never above a fact that organization has already confirmed.

Chapter 7 — Invisible Automation

The best automation is the one users barely notice.

Routine work happens automatically.

Learning happens continuously.

Notifications are meaningful.

Every interruption must have value.

Success is measured by how much work Centro completes without demanding attention.

Chapter 8 — The Boundary of Learning

Centro learns exactly one thing: **which documents should be requested from each client.**

This is refined only through observed recurrence and explicit client confirmation, following the Chapter 3 lifecycle exactly — never inferred, never assumed, and never applied from a single occurrence.

Centro does not learn, infer, or adapt:

- The client's underlying business, workflow, or operations.
- The accounting firm's own office policies — including business hours, working days, reminder cadence, or the day a collection period begins.

These remain office policy: configured explicitly by the accountant, once, and changed only by the accountant. Centro never modifies them automatically, regardless of what it observes.

This boundary is permanent. It applies to every learning mechanism in Centro — present and future — not only document collection. A future milestone that would cause Centro to learn or adapt timing, scheduling, or any office policy is out of scope by definition, not a case-by-case judgment call.

Extension (Product Evolution M7): the boundary now has a second, absolute form. For a `one_time` - workflow organization (Chapter 9), Centro does not learn *at all* — not a client's document profile, not a recurring pattern, nothing. This is stronger than "refined only through observed recurrence": it is "never even observed." The guarantee is enforced inside the three functions that ever write learned state (`recordAdHocDocumentObservation`, `detectMissingRequirements`, `recordLearnedDocumentPattern`), each checking `isOneTimeWorkflowOrganization(organizationId)` as their first line — not only by Workflow B's UI never exposing the paths that would trigger it. A future caller that forgets this check still cannot cause Workflow B to learn.

Chapter 9 — The Two-Workflow Product Model

Everything in Chapters 1–8 describes Centro as built for one kind of business: recurring, document-collecting professional services (the accounting-firm flagship). The Product Evolution roadmap generalized

Centro into a platform for *any* document-collecting business, without rewriting that flagship experience. It did this by introducing exactly one new, permanent fork, decided once at onboarding and never silently changed afterward.

9.1 Two Facts, Set Once

Every organization declares two things during onboarding that no chapter before this one assumed varied:

- `organizations.businessCategory` — what kind of business the office itself is (accountant, tax advisor, lawyer, real estate, mortgage advisor, insurance, HR, finance, or a free-text "other"). This is deliberately distinct from the pre-existing `business_types` table, which classifies the office's *clients* and is a Workflow-A-only concept. Category is used for onboarding personalization and AI-driven document suggestions — never to change how the learning engine itself behaves.
- `organizations.workflowType` — `recurring` (Chapters 1–8, exactly as written) or `one_time` (this chapter). Every organization that existed before this roadmap defaults to `recurring`, so nothing in Chapters 1–8 changed for a single existing organization the day this shipped.

Both are meant to be permanent, set-once onboarding decisions — not a Settings toggle a user can flip later. This is enforced structurally, not just by omission: once `organizations.onboardingCompletedAt` is set, the onboarding wizard itself becomes unreachable (Product Evolution M8) — a completed organization visiting `/onboarding` directly is redirected straight back to its dashboard before any onboarding action, including a workflow-type change, can run. Switching workflows remains possible only as a deliberate, out-of-band data-migration operation, never as a page a logged-in user can stumble back into.

9.2 Workflow A (Recurring) — Unchanged, Now Generalized at the Edges

For an `accountant` or `tax_advisor` organization, Workflow A is byte-identical to Chapters 1–8: the same wizard, the same Excel analysis, the same starter business types, the same learning engine. For every other declared category, only the onboarding *starting point* generalizes — a mocked-AI module (`src/lib/ai/businessCategorySuggestions.ts`) supplies category-appropriate starter client types and document checklists instead of the five hardcoded Israeli accounting types, falling back to keyword-matching and then a broad, never-empty "General Client" starter for a genuinely novel free-text category. Once seeded, a business type behaves identically regardless of which tier produced it — Chapter 3's lifecycle, Chapter 6's Decision Hierarchy, and the learning engine itself never change based on category. Only the starting point does.

9.3 Workflow B (One-Time) — A Different, Smaller, Non-Learning Product

A `one_time`-workflow organization gets a materially different experience from the moment onboarding forks (after the shared Welcome → Office Info → Business Type → Collection Style steps): an optional client import that stores only name and phone (no classification, no analysis), then Working Hours only (no reminder cadence or collection-day-of-month — office-policy concepts with no meaning for a one-off request), then its own summary and completion screen. It lands on its own dashboard and navigation, built around one new primary surface:

A Template is a bare `services` row belonging to a `one_time`-workflow organization — no new table. A Template has a name, a list of required documents (AI-suggested per business category via `suggestTemplateLibrary`, fully user-editable — add, rename, reorder, remove), and any number of assigned clients (`client_services`, the same join table Workflow A's service assignment already used). Sending a Template creates one ordinary `collectionRequest` per assigned client — reusing `snapshotServiceRequirements` unchanged — and delivers it either immediately or at a scheduled future time via the single shared `attemptScheduledDelivery` function (`src/lib/scheduledSend.ts`); there is no separate "send now" code path. From the moment of delivery onward, Workflow B's collection requests flow through the *exact same* shared infrastructure as Workflow A's: WhatsApp messaging, OCR/classification (Chapter 6), Google Drive's one-folder-per-client organization, document tracking, and the reminder engine (Chapter 7) — all reused without modification.

The one thing Workflow B never does, by design, is learn (see Chapter 8's extension above). Templates are entirely explicit and user-edited; there is no per-client document-profile history to observe, because a one-time office's relationship with a client is, definitionally, one-time.

9.4 The Permanent Distinction, In One Sentence

Recurring Workflow: Centro learns document requirements over time, per client. **One-Time Workflow:** Centro never learns; the user manages reusable Templates instead. Every other product decision in this chapter follows from that one line.

9.5 UX Polish — Terminology & Surfaces

A dedicated UX-only round (no architecture, no business logic) brought the product's own terminology in line with the two-workflow model this chapter describes, and closed two onboarding-era gaps:

- **"Office" → "Business."** The office's own profile (name, logo, the Settings section that edits them) now reads "Business" throughout — the office/firm-specific framing didn't fit a platform meant for lawyers, HR teams, and mortgage advisors as much as it fit the accounting flagship. Client-facing automated copy (the WhatsApp greeting) and public marketing pages were deliberately left untouched — different audience, out of scope for an internal-app terminology pass.
- **"Services" → "Templates," extended to Workflow A's own surface.** §9.3 already established Templates as Workflow B's vocabulary for a bare `services` row. This round extended the same word to Workflow A's own `/services` pages (route, table, and every internal name unchanged — text only), plus the two pages genuinely shared by both workflows (a client's assigned-services section, `/collections`), which now render "Service" or "Template" based on the viewing organization's own `workflowType`. This also closed a real pre-existing gap: `/services/*` had no workflow gate at all, unlike `/templates/*`'s existing one, meaning a one-time organization could reach the Collection Day override field this same round hides everywhere else. `/services/*` now 404s for `one_time` organizations, symmetric with `/templates/*`.
- **"Audit Log" → "Activity History,"** with a real behavior change to match: the default view is now today's events (not the last 200 rows, unbounded by date), with Today/Last-7-Days/Last-30-Days/Custom-Range filters and day-grouped results. `listAuditLog`'s new `from / to` bounds are opt-

in — every other caller (e.g. a Collection Request's own full history on `/collections/[id]`) is unaffected.

- **Two new small state flags**, both following this document's established idiom (a boolean on the owning row for a permanent fact, a nullable timestamp on `organizations` for "has this one-time thing happened yet"): `services.isSampleTemplate` (true only for the templates `seedExampleTemplates` auto-creates, never for a user's own library addition or duplicate — see §9.3) powers a "Sample Template" badge and a short explanation banner; `organizations.sampleTemplateCardShownAt` gates a one-time dashboard promo card to exactly one appearance, the organization's first visit.
- **A complete password-reset flow** (`password_reset_tokens`) was added alongside login UX polish (show/hide password, a dual-duration "Remember Me," a real `/register` URL). The reset-email step is mocked per this document's established pattern for pilot-stage external dependencies with no live provider configured (Chapter 6's classifiers are the same shape of stand-in) — the reset link is logged server-side, never exposed in any HTTP response, and the flow never reveals whether a given email is registered.

None of this round touched Chapters 1–9's actual product logic — every change here is a label, a default filter, a small new flag, or closing a routing gap the two-workflow model itself already implied should exist.

Document History

Date	Change
Pilot Edition	Original 7 chapters established as the architectural source of truth.
Milestone 1	Migrated to a version-controlled Markdown source in-repo. Chapter 6 revised to state the Decision Hierarchy as actually implemented (Learned Knowledge checked before generic Business Rules) — see the Centro Implementation Report's Milestone 1 addendum for the full rationale. Added Chapter 8 (The Boundary of Learning) as a permanent constraint, in response to an early roadmap draft that would have violated it.
Milestone 2	Chapter 1/2's "every client begins in Learning Mode" implemented as a real, visible flag (<code>clients.learningMode</code>), ending the first time a client's first Collection Request reaches <code>completed</code> . No timing or duration data is captured — see Chapter 8; this milestone exists solely to gate Milestone 6's "document appears to no longer be needed" detection, which must never fire before a client has completed even one full cycle.
Milestone 3	Chapter 6's Decision Hierarchy generalized to document classification, its second real implementation (after Epic 4's business-type classifier). A client's own manually-corrected document-naming history (<code>document_learned_patterns</code> , scoped per client) is now checked before the generic deterministic heuristic, which now sits ahead of a documented AI-fallback stub — full parity with the 4-layer pattern.
Milestone 4	Chapter 4/7's "the accountant reviews exceptions, not routine learning" given a real, unified surface: the dashboard's existing per-feature queue pattern extended with a "suggested classification" view (clients auto-classified in the 70-94% band), sitting alongside the pre-existing document-review queue rather than a separate, undiscoverable screen.

Date	Change
Milestone 5	Chapter 3's "Confirm with the client, through WhatsApp" — the one lifecycle step that didn't exist anywhere yet — built as reusable, domain-agnostic infrastructure (<code>pending_confirmations</code>), with a real (not throwaway) manual trigger: an employee can ask a client to confirm a document becomes a standing part of their collection at the exact moment of manual assignment. A real bug was found and fixed during this milestone's own verification — see the Implementation Report's Milestone 5 addendum.
Milestone 6	Chapter 8's one permitted kind of learning, made real end to end: <code>client_document_requirements</code> tracks per-client additions and removals, confirmed only through Milestone 5's WhatsApp loop, applied only to that one client's future snapshots — the service template itself is never touched. Additions require a genuine second occurrence (never a single event); removals require two consecutive completed cycles missing the same document. A structural constraint (a cycle can never complete while a requirement is unsatisfied) required a new, independently useful capability — waiving a requirement for one cycle — to make recurring-absence detection possible at all; documented in the Implementation Report's Milestone 6 addendum along with a second bug (ambiguous reply resolution when two confirmations are open at once) found and fixed during verification.
Milestone 7	Pilot Hardening — no new product behaviour; a dedicated pass over everything M1–M6 built, with an emphasis on the organization-isolation boundary this document has assumed throughout but never itself audited end to end. Two real cross-tenant authorization gaps were found and fixed (a pending-confirmation resolution action, and a document-requirement-assignment action, both trusting a caller-supplied id without confirming it belonged to the caller's own organization) — see the Implementation Report's Milestone 7 addendum for the full detail, including a two-tenant Playwright test added specifically to keep this boundary verified going forward, not just asserted in prose. The Milestone 3-deferred <code>isFuzzyDuplicate</code> doc-comment discrepancy was also resolved (the comment's own example didn't clear its own threshold — corrected to a real, passing example; the function's deliberately strict behavior was left unchanged).
Product Evolution M1	The start of a separate roadmap generalizing Centro from an accounting-only platform into a document-collection platform for any business (see the Product Evolution milestone roadmap). This milestone captures two new, purely onboarding-time facts and changes no existing behaviour: which kind of business the office itself is (<code>organizations.businessCategory</code> — distinct from <code>business_types</code> / <code>clients.businessTypeId</code> , which classify this office's <i>clients</i> and remain exactly as Chapters 6 and 8 describe), and which of two permanent workflows it operates in (<code>organizations.workflowType</code> : <code>recurring</code> , today's exact learning-driven experience, or <code>one_time</code> , a template-driven experience with deliberately no learning at all, built out in later milestones). Every existing organization defaults to <code>accountant / recurring</code> , so nothing in Chapters 1–8 changes for any organization that existed before this milestone. The one-time workflow's own governing chapter will be added once it's real (the Product Evolution roadmap's final hardening milestone).
Product Evolution M2	Generalizes Workflow A's onboarding <i>starting point</i> — never the recurring learning engine itself — per the office's declared business category. For <code>accountant / tax_advisor</code> , the starter client-types and suggested documents are byte-identical to before this milestone. For every

Date	Change
	other declared category, a new mocked-AI module (<code>src/lib/ai/businessCategorySuggestions.ts</code>, same documented "real interface, no live provider configured" pattern as every other AI module in this codebase) supplies curated, category-specific starter types — and, for a genuinely novel free-text "Other" category, a broad keyword-matching layer covering ~20 common business domains before falling back to a starting point named after the office's own words, never a generic or empty placeholder. Once seeded, a business type behaves identically regardless of which tier produced it: Chapter 3's Observe → Suggest → Confirm → Learn lifecycle, Chapter 6's Decision Hierarchy, and the entire recurring learning engine apply exactly as before — this milestone only ever changes what a client-type's document checklist starts as, never how Centro learns from that point on.
Product Evolution M3	The onboarding wizard genuinely forks for the first time. Steps 1–5 (Welcome through Connect) stay identical for both workflows; from Step 6 onward, <code>organizations.workflowType</code> determines a completely different continuation. Workflow A's Steps 6–11 are untouched (regression-verified). Workflow B gets its own, deliberately shorter path (9 steps total): an explicitly optional Client Import step that stores <i>only</i> name and phone — no classification, no learned-synonym writes, no <code>business_types</code> seeding, no analysis summary, honoring the product's "Centro never learns" boundary for this workflow from the very first cycle, not just later once Templates exist — followed by a Working Hours step (business days/hours only, deliberately never asking about reminder cadence or collection-day-of-month, which remain office-policy concepts with no meaning for a one-off request), then an adapted Summary and Completion. <code>WizardShell</code>'s step count and titles are now resolved per request instead of fixed constants, the deferred piece of work flagged at the end of Milestone 1. A full live walk of both paths (recurring regression and both one-time branches) passed cleanly on the first run — the step-renumbering bug class Milestones 1 and 2 both found instances of had already been fully swept by Milestone 2's codebase-wide check.
Product Evolution M4	One-time-workflow organizations get their own dashboard and navigation, sharing the underlying route and shell rather than a separate layout — matching the spec's own "Shared Infrastructure" list, which names Dashboard Infrastructure as shared. <code>/dashboard</code> branches at the top by <code>organizations.workflowType</code>; the sidebar's nav list does the same. The one-time dashboard reuses the recurring dashboard's <i>pattern</i> (small, independent, composable query functions) but none of its queue vocabulary, since business-type suggestions, pending classification confirmations, and the rest are Workflow-A-only concepts. A minimal <code>/templates</code> placeholder, guarded to one-time organizations only, keeps the new "תבניות" nav link and dashboard CTA from being dead links until Milestone 5 builds the real Templates CRUD. A recurring organization's dashboard and navigation are byte-identical to before this milestone, confirmed via live regression.
Product Evolution M5	The real Templates CRUD, replacing Milestone 4's placeholder — create/edit/delete/duplicate a Template, and add/rename/reorder/remove its document requirements, all as thin Template-branded wrappers around the exact operations <code>/services</code> already had (a Template is a bare <code>services</code> row). Reordering is simple move-up/move-down rather than drag-and-drop, deliberately — new <code>service_document_requirements.position</code> column, null for every pre-existing row and for the recurring workflow generally, so ordering falls back to creation order with no special-case code. The Template Library (<code>suggestTemplateLibrary</code>, extending

Date	Change
	Milestone 2's AI-suggestions module) offers one-click starter templates per business category, including a dedicated set for <code>accountant / tax_advisor</code> distinct from Workflow A's client-classification starters — a one-time office's request templates and a recurring office's client-type taxonomy are different concepts even for the same declared category. The first visit to <code>/templates</code> auto-seeds three Library suggestions as real, immediately editable templates, mirroring <code>seedStarterBusinessTypes</code> 's own idempotent-on-emptiness pattern.
Product Evolution M6	Client assignment for Templates — reuses <code>client_services</code> directly (the same join table the existing client-detail-page "assign service" action already writes) rather than a new table, since "clients assigned to this template" and "clients assigned to this service" are the same relationship once a Template is understood as a bare <code>services</code> row. The Template detail page gained an assignment panel mirroring the onboarding wizard's own unclassified-clients picker (a checkbox list + submit), plus an inline "create new client" shortcut for the common one-time-workflow case of assigning a template to someone not in the system yet. A client can be assigned to any number of templates simultaneously, and the assignment panel deliberately stays open after a successful assignment so several clients can be added in one sitting.
Product Evolution M7	Send Request — the step that actually delivers a Template to its assigned clients, closing the loop into the existing collection pipeline. One submit creates a <code>collectionRequest</code> per selected client (reusing <code>snapshotServiceRequirements</code> , unchanged) and either sends immediately or schedules a future delivery — deliberately one mechanism, not two: a new nullable <code>collection_requests.scheduledAt</code> is null for "send now" (delivered synchronously) and non-null for "schedule for later," and the exact same function, <code>attemptScheduledDelivery</code> , delivers both — the only difference is who calls it and when. The cron tick (<code>runScheduledTasks</code>) gained a due-scheduled-request query so a "send now" that lands outside business hours degrades gracefully into "delivered on the next tick" rather than silently failing, using infrastructure that already existed for reminders. Chapter 8's learning boundary is now enforced at the three functions that ever write learned state (<code>recordAdHocDocumentObservation</code> , <code>detectMissingRequirements</code> , <code>recordLearnedDocumentPattern</code>), not only by omission in Workflow B's UI — a new <code>isOneTimeWorkflowOrganization</code> check is the first line of each, so the guarantee holds even if a future caller forgets to check first. Verified live, including the full scheduled-delivery path (a request scheduled ~65s out was confirmed to stay <code>draft</code> until its time arrived, then flip to <code>active</code> via a real cron-tick run) and the learning guard (the exact two-occurrence condition that triggers a client-confirmation prompt for a recurring client was reproduced on a one-time client and confirmed to never surface it).
Product Evolution M8	Pilot Hardening for the whole epic — added Chapter 9 (The Two-Workflow Product Model) and extended Chapter 8 with Workflow B's absolute "never observed" form of the learning boundary; no prior chapter's content changed. A cross-tenant Playwright audit (two independent one-time organizations) confirmed every new M1–M7 surface — <code>/templates</code> , <code>/templates/[id]</code> , the one-time dashboard, <code>sendTemplateRequest</code> , and every Template CRUD/assignment action — correctly rejects or hides another organization's data, with zero new isolation gaps found; every action was already scoping by <code>session.organizationId</code> , a direct consequence of M5–M7 building each one as a thin wrapper around the pre-existing,

Date	Change
	already-audited <code>services / clients</code> data-access functions rather than new parallel ones. One genuine gap <i>was</i> found and fixed: nothing previously stopped a fully onboarded organization from navigating straight back to <code>/onboarding</code> and resubmitting <code>updateWorkflowType</code> , silently flipping a live organization between the recurring and one-time workflows post-onboarding — contradicting this epic's own "permanent, set-once" design (§9.1). Fixed with one guard, mirroring the <code>(app)</code> layout's existing reverse-direction check: once <code>onboardingCompletedAt</code> is set, <code>/onboarding</code> itself redirects to <code>/dashboard</code> before any onboarding action can run. A second, smaller gap — four of <code>sendTemplateRequest</code> 's rejection paths (invalid schedule, no clients selected, etc.) redirected with an <code>?error=</code> code the template page never actually displayed, leaving the user with silent, unexplained failures — was also found and fixed. Additional edge cases verified without incident: a ~1,000-character free-text business category produces sensible, non-crashing Template Library suggestions; a template with zero assigned clients simply hides the Send Request panel rather than presenting a dead-end form.
UX Polish M1–M8	A dedicated UX-only pass (Chapter 9, §9.5) — no architecture or business-logic change. Terminology aligned with the two-workflow model ("Office"→"Business," "Services"→"Templates" for Workflow A's own pages, "Audit Log"→"Activity History" with a real default-to-today rewrite), plus a <code>/services/*</code> workflow gate closing a pre-existing gap symmetric with <code>/templates/*</code> 's own. New <code>services.isSampleTemplate / organizations.sampleTemplateCardShownAt</code> flags power a one-time "Sample Template" onboarding nudge. A complete password-reset flow (<code>password_reset_tokens</code>) shipped alongside login UX polish (show/hide password, a dual-duration "Remember Me," a real <code>/register</code> URL). One real, pre-existing bug was found and fixed during this round's own regression pass: <code>TemplateSendRequest</code> 's client-selection state was seeded once from <code>useState</code> 's initializer and never updated when a client was assigned <i>after</i> the component first rendered, leaving "Send" permanently disabled at 0 selected until a full page reload — fixed by adjusting the selection during render when the <code>assignedClients</code> prop actually changes (React's own documented pattern for this case), not via a <code>useEffect</code> .
UI Redesign M0–M9	A full visual and interaction redesign of the authenticated application (login through every internal screen), executed against the Centro Design Manifesto v1.0. Zero product-logic change — no schema, API, workflow, or business-rule modification; every chapter above remains exactly as implemented. Purely additive Tailwind design tokens (<code>src/app/globals.css</code>), a new shared component library (<code>src/components/app/</code> : <code>Table</code> , <code>Tabs</code> , <code>KpiCard</code> , <code>AuthCard</code> , <code>ConfirmDialog</code> , <code>DevToolsPanel</code> , <code>ProgressBar</code> , <code>AiBriefing</code> , an extended <code>FormField</code>), and a per-screen visual migration across the app shell, onboarding wizard, both dashboards, clients, services/templates, collections, settings, and activity history — the latter two's zero-client-JS, bookmarkable-URL filter mechanisms preserved exactly, restyled only. The public marketing Landing Page was explicitly out of scope and confirmed pixel-identical throughout (zero diff on <code>src/components/landing/**</code> across all ten milestones). Two real cross-cutting bugs were found and fixed during this epic's own verification, both documented in the redesign's own final report: <code>ConfirmDialog</code> 's original <code>cloneElement</code> -based trigger API broke intermittently under React Server Components (fixed by accepting plain <code>ReactNode</code> content instead of an element to clone), and this document's own Architecture PDF

Date	Change
	(<code>reports/Centro-Architecture-and-Philosophy.pdf</code>) was discovered to be a broken artifact — a printed "file not found" browser error page, not the document's actual content — regenerated correctly as part of this pass.
UI Redesign — Centro Glow Phase 2	A second, iterative visual round on top of M0–M9's redesign, refining it into the "Centro Glow" language: near-white surfaces with extremely subtle per-card hover glow (behind the card, color-matched to its own icon, never inside it), soft gradient icon badges, and a small "Centro Active" status indicator on the Dashboard only (a breathing dot + tooltip, aligned to the page title via real flex layout, not a guessed offset). The onboarding wizard got the most visible changes: the official <code>CentroMark</code> brand mark now appears at the top of every step and on Login/Register (glow/breathe wrapper only — the mark's own SVG geometry is never touched, a permanent rule); the Welcome step's animated hand was removed entirely in favor of a staggered typography-and-motion reveal (glow → eyebrow → blur-to-sharp title → subtitle → rising button) after review found every CSS/SVG attempt at a hand read as artificial; Service Connections shows real Google Drive/WhatsApp brand marks and a shared circle-draws-into-checkmark animation (<code>AnimatedCheckBadge</code>); the AI-analysis step reveals progressively (title, card, text, then each stat counting up) instead of appearing instantly; Setup Summary rows animate in with the same checkmark draw, staggered; and onboarding completion replaced a static particle badge with a real canvas confetti burst (blue/purple/teal/white/gold, ~1.7s). Zero product-logic change — every refinement is visual/motion only, built from static mockups reviewed and explicitly approved screen-by-screen before implementation. One real bug was found and fixed during this round's own live verification: two CSS classes applied to the same Setup Summary row both set the <code>animation</code> shorthand, and the later rule silently discarded the earlier one's <code>animation-name</code> entirely (equal specificity) — every "done" row was stuck invisible until the two animations were merged into one shorthand declaration.
Product Fixes & Mobile Responsiveness	A focused round of product fixes plus a full-product mobile responsiveness audit, no visual redesign. Fixes: real, complete Hebrew Privacy Policy and Terms of Service content (replacing placeholder text), now linked from Login/Register and every onboarding step's footer, not only the registration checkboxes. The one-time workflow's client import was brought to parity with the recurring workflow's upload/replace/add-another controls — reusing the exact same <code>clients.importedDuringOnboarding</code> flag and replace-deletion query the recurring import already relies on (no schema change; the flag already existed as a generic column, just previously unset by <code>importClientsSimple</code>). The import step now stays in place after a successful upload to show a review state with Replace/Add-another and an explicit Continue, instead of auto-advancing — mirroring the self-redirect-to-the-same-step trick <code>importParsedRows</code> already used for the recurring workflow's own review step. Setup Summary rows now show a red X for an incomplete/skipped step instead of an empty hollow circle, in both workflows, row text unchanged. The one-time Dashboard's four KPI cards are now clickable (Templates, Clients, Collection Requests), reusing <code>KpiCard</code> 's existing <code>href</code> behavior — zero visual change. Internal reference codes in parentheses (e.g. (BR-18.1) , (Ch.17)) were removed from the user-facing Settings and Activity History descriptions. Responsiveness: a systematic audit across the landing page, every auth screen, the full onboarding wizard (both workflows), and every main app screen at five phone/tablet viewport widths found one real bug — the Welcome screen's ambient glow blob paired a logical Tailwind position class (<code>start-1/2</code>) with a physical transform (<code>-translate-x-1/2</code>); correct in LTR, but under this app's RTL

Date	Change
	layout the two no longer cancel out, and the blob was pushed far outside the viewport, forcing real horizontal overflow on every phone size (fixed with the matching physical <code>left-1/2</code>). Every other flagged page turned out to be a <code>scrollWidth</code> measurement artifact already neutralized by the existing <code>body { overflow-x: clip }</code> safeguard, confirmed via a direct <code>scrollTo</code> -based check rather than assumed.
Landing Page — Contact Form Email & WhatsApp Button	Two additions to the public landing page only; the app's authenticated surfaces and the landing page's own visual design were untouched. Contact form → email: <code>ContactForm.tsx</code> (shared by <code>ContactSection</code> and <code>DemoRequestModal</code>) previously simulated a network call and always "succeeded." It now posts to a new <code>POST /api/contact</code> route handler, which re-validates the payload server-side (never trusting client-side validation as a security boundary) and sends a formatted HTML+text email via Resend to a configurable inbox (<code>CONTACT_EMAIL_TO</code> , defaulting to <code>Centro.ai.team@gmail.com</code> in code) from a configurable sender (<code>RESEND_FROM_EMAIL</code> , defaulting to Resend's sandbox sender), including every submitted field and the submission timestamp. The success UI only ever renders after a real <code>2xx</code> from the route; any failure (missing <code>RESEND_API_KEY</code> , a Resend-side error, a network failure) surfaces as an inline error message and re-enables the submit button for retry, rather than a false-positive success. The submit button already guarded against duplicate submission via its <code>status === "submitting"</code> disabled state; a redundant top-of-handler guard was added for defense in depth. <code>RESEND_API_KEY</code> is read only in the route handler (server-side), never sent to the client. Floating WhatsApp button: a new <code>FloatingWhatsAppButton.tsx</code> , mounted only from <code>src/app/page.tsx</code> (never the root layout), so it appears exclusively on the landing page and never on any authenticated or auth-flow screen. Fixed bottom-right (physical <code>right-*</code> , not the logical <code>end-*</code> that would resolve to the opposite corner under this app's RTL layout — the same class of mistake Mobile Responsiveness's one real bug above was), respecting <code>env(safe-area-inset-*)</code> , positioned opposite <code>AccessibilityWidget</code> 's existing bottom-left placement so the two never collide. Reuses the existing <code>WhatsAppGlyph</code> icon and <code>--color-whatsapp</code> design token rather than introducing new brand colors, with one new additive keyframe (<code>centro-whatsapp-breathe</code> , a slow ambient glow pulse behind the button, echoing the app's established "glow behind, never inside" hover mechanic) — already covered by the existing global and manual <code>prefers-reduced-motion</code> overrides with no bespoke exception needed. Opens <code>https://wa.me/972559871812</code> with a pre-filled Hebrew message, <code>target="_blank" rel="noopener noreferrer"</code> , and a Hebrew <code>aria-label</code> for screen readers.